

HTML HW2

1. A hypothesis fixed at point can rotate freely at any angle. Now, to separate the points, we must place the line between the ones with output 1 and the ones with output -1 . There are $N + 1$ possibilities, but the output 1 and -1 is interchangeable, so the answer is $2N + 2$, [c].
2. For a d -dimensional perceptron, the VC dimension is $d_{VC} = d + 1$, and with x_0 fixed, $d_{VC} = 6211$. But there is only 1126 hypotheses, so the tighter upper bound is $\log_2(1126)$, [a].
3. [a] For the union of two positive interval, the growth function is of order

$$m_{\mathcal{H}}(N) \sim \binom{N+1}{4} = O(N^4)$$

but according to the theorem

$$m_{\mathcal{H}}(N) = O(N^{d_{VC}})$$

we conclude that $d_{VC} = 4$. For $N = 5$, the dichotomy OXOXO can not be separated.

[b] For d -dimensional perceptrons, $d_{VC} = d + 1 = 4$.

[c] The function $x - \sin(\omega x)$ oscillates between positive and negative more frequently with larger ω . As $\omega \rightarrow \infty$, the function changes sign approximately every π/ω , or infinitely often, which means with a suitable choice of ω , and the inputs constrained not far from the origin, we can find some N inputs that can be shattered by \mathcal{H} , which means $d_{VC} \rightarrow \infty$.

[d] For $N = 1, 2$, it can be easily shown that the inputs can be shattered. For $N = 3$, position the inputs such that the points form a right triangle with the right angle in the top left corner. This way we can shatter all inputs. For $N = 4$, position the points as a diamond.

O

O O

O

then it can be shown that the inputs can be shattered. For $N = 5$, place a negative point in the interior of the other four positive points, then it is not possible to shatter the inputs, so $d_{VC} = 4$.

[e] For $N = 3$, position the data points as follows

O O

O

then the inputs can be shattered. For $N = 4$, we can not shatter the inputs no matter how we place the points, so $d_{VC} = 3$.

The answer is [e].

4.

Sort each input by its value of $\mathbf{x}^T \mathbf{x}$, and have its label alternate between $+1$ and -1 for $2m + 1$ inputs. Then the points can not be shattered, which means $d_{VC} \leq 2m$. To prove $d_{VC} \geq 2m$, we must prove that we can shatter some $2m$ inputs. We know that the value of the hypothesis alternates as $- + - + \dots + -$ for $2m + 1$ times. We again choose inputs that alternate signs every time, because consecutive $+$ or $-$ can be grouped into the same interval and the input can be regarded as alternating every time but with a shorter length. So, the choices are $- + - + \dots - +$ or $+ - + - \dots + -$ with length $2m$, both of which can be shattered by the hypothesis $- + - + \dots + -$, and inputs with consecutive signs is equivalent to $- + - + \dots$ with length $< 2m$ which can also be shattered. so we have proved $d_{VC} \leq 2m$.

The answer is [b].

5. Having a VC dimension of d means that we can not find any set of $N > d$ inputs that can be shattered. But we can find some d inputs we can shatter. So $d_{VC} \leq d$ implies that any set of $d + 1$ distinct inputs is not shattered. which also implies that some set of $d + 1$ distinct inputs is not shattered. Condition 6 and 8 is correct, the answer is [b].

6. We would like to minimize $E_{in}(w)$ by taking its derivative

$$\frac{dE_{in}(w)}{dw} = \frac{2}{N} \sum_{n=1}^N x_n (wx_n - y_n) = 0$$

It follows that

$$w = \frac{\sum_{n=1}^N x_n y_n}{\sum_{n=1}^N x_n^2}$$

The answer is [b].

7. We would like to choose the probability distribution that is most likely to produce the given data set, which has a probability of

$$P = P(x_1) \times P(x_2) \times \dots \times P(x_N)$$

We vary the parameter of the distribution to find the argument for maximum P .

[a] For a Poisson distribution,

$$P = \frac{e^{-N\lambda} \lambda^{x_1+x_2+\dots+x_N}}{x_1! x_2! \dots x_N!}$$

$$\frac{dP}{d\lambda} = \frac{1}{x_1! x_2! \dots x_N!} (-N e^{-\lambda} \lambda^{\sum_n x_n} + e^{-\lambda} \lambda^{-1+\sum_n x_n} \sum_{n=1}^N x_n) = 0$$

Reorganizing

$$\lambda = \frac{1}{N} \sum_{n=1}^N x_n$$

[b]

$$P = (2\pi)^{-N/2} e^{-\frac{1}{2} \sum_n (x_n - \mu)^2}$$

$$\ln(P) = -\frac{N}{2} \ln(2\pi) - \frac{1}{2} \sum_{n=1}^N (x_n - \mu)^2$$

$$\frac{d \ln P}{d\mu} = \sum_{n=1}^N (x_n - \mu) = 0$$

It follows that

$$\bar{x} = \mu$$

[c]

$$P = 2^{-N} e^{-\sum_n |x_n - \mu|}$$

We would like to find the value of μ that minimizes, sort

$$\sum_{n=1}^N |x_n - \mu| = |x_1 - \mu| + |x_2 - \mu| + \cdots + |x_N - \mu|$$

sort x_n by ascending order, and use the inequality

$$|a - b| + |c - b| \geq |a - c|$$

we have

$$\sum_{n=1}^N |x_n - \mu| \geq |x_N - x_1| + |x_{N-1} - x_2| + \cdots$$

The condition for the equal sign to hold is that μ must be the median, which is not always equal to the mean \bar{x} . Therefore this claim is incorrect.

[d]

$$P = (1 - \theta)^{x_1 + x_2 + \cdots + x_N - N} \theta^N$$

$$\frac{d \ln(P)}{d\theta} = \frac{d}{d\theta} ((x_1 + x_2 + \cdots + x_N - N) \ln(1 - \theta) + N \ln \theta) = \frac{N}{\theta} - \frac{-N + \sum_n x_n}{1 - \theta} = 0$$

solving for θ

$$\frac{1}{\theta} = \frac{1}{N} \sum_{n=1}^N x_n$$

The answer is [c].

8. since $\tilde{h}(x) = \tilde{h}(-x)$,

$$P(\mathbf{x}_1)h(\mathbf{x}_1) \times P(\mathbf{x}_2)h(-\mathbf{x}_2) \times \cdots \times P(\mathbf{x}_N)h(-\mathbf{x}_N) \propto \prod_{n=1}^N h(y_n \mathbf{x}_n)$$

$$\max_{\mathbf{w}} \ln \prod_{n=1}^N h(y_n \mathbf{x}_n) \Rightarrow \min_{\mathbf{w}} \sum_{i=1}^N -\ln h(y_n \mathbf{x}_n)$$

Therefore

$$\tilde{E}_{\text{in}}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \ln \left(\frac{2 + 2|y_n \mathbf{w}^T \mathbf{x}_n|}{1 + y_n \mathbf{w}^T \mathbf{x}_n + |y_n \mathbf{w}^T \mathbf{x}_n|} \right)$$

Taking its gradient

$$\frac{\partial \tilde{E}_{\text{in}}(\mathbf{w})}{\partial w_i} = -\frac{1}{N} \sum_{n=1}^N \frac{y_n x_{ni}}{(|y_n \mathbf{w}^T \mathbf{x}_n| + 1)(|y_n \mathbf{w}^T \mathbf{x}_n| + y_n \mathbf{w}^T \mathbf{x}_n + 1)}$$

thus

$$\nabla \tilde{E}_{\text{in}}(\mathbf{w}) = -\frac{1}{N} \sum_{n=1}^N \frac{y_n \mathbf{x}_n}{(|y_n \mathbf{w}^T \mathbf{x}_n| + 1)(|y_n \mathbf{w}^T \mathbf{x}_n| + y_n \mathbf{w}^T \mathbf{x}_n + 1)}$$

The answer is [a].

9.

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

$$\nabla E_{\text{in}}(\mathbf{w}) = \frac{2}{N} (\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y})$$

The definition of the Hessian matrix is

$$(\nabla^2 E_{\text{in}}(\mathbf{w}))_{ij} = \frac{\partial^2 E_{\text{in}}(\mathbf{w})}{\partial w_i \partial w_j}$$

let $\mathbf{A} = \mathbf{X}^T \mathbf{X}$ and $\mathbf{b} = \mathbf{X}^T \mathbf{y}$, since

$$\frac{\partial E_{\text{in}}(\mathbf{w})}{\partial w_i} = \frac{2}{N} \left(\sum_{j=1}^N \mathbf{A}_{ij} \mathbf{w}_j - b_i \right)$$

then

$$(\nabla^2 E_{\text{in}}(\mathbf{w}))_{ij} = \frac{\partial^2 E_{\text{in}}(\mathbf{w})}{\partial w_i \partial w_j} = \frac{2}{N} A_{ij}$$

thus

$$\nabla^2 E_{\text{in}}(\mathbf{w}) = \frac{2}{N} \mathbf{X}^T \mathbf{X}$$

The answer is [b].

10.

$$(\nabla^2 E_{\text{in}}(\mathbf{w}_t))^{-1} = \frac{N}{2} \mathbf{X}^{-1} (\mathbf{X}^T)^{-1}$$

$$\mathbf{w}_1 = \mathbf{u} = -\frac{N}{2} \mathbf{X}^{-1} (\mathbf{X}^T)^{-1} \left(-\frac{2}{N} \mathbf{X}^T \mathbf{y} \right) = \mathbf{X}^{-1} \mathbf{y}$$

If we check the gradient

$$\nabla E_{\text{in}}(\mathbf{w}_1) = \frac{2}{N} (\mathbf{X}^T \mathbf{X} (\mathbf{X}^{-1} \mathbf{Y}) - \mathbf{X}^T \mathbf{y}) = \mathbf{0}$$

The answer is [a].

11. The VC bound is given by

$$P_{\mathcal{D}}[|E_{\text{in}} - E_{\text{out}}| > \epsilon] \leq 4(2N)^{d_{\text{vc}}} \exp(-\epsilon^2 N/8)$$

For the decision stump model

$$\delta = 16N^2 \exp(-\epsilon^2 N/8)$$

Calculating δ for each option

[a] 1.55×10^5

[b] 1.17×10^7

[c] 7.03×10^7

[d] 4.29×10^{-3}

[e] 3.07×10^{-123}

The answer is [d].

12. There are multiple scenarios depending on the range of θ and x :

- $\theta > 0.5$: Here $h(x) = -1$, therefore

$$E_{\text{out}} = P[y = +1] = \tau P(x < 0) + (1 - \tau) P(x > 0) = \frac{1}{2}$$

- $0 < \theta < 0.5$: Consider three cases,

1. $\theta < x \leq 0.5$: Here $h(x) = +1$, thus

$$E_{\text{out}} = P[y = -1] = \tau$$

2. $0 < x \leq \theta$: $h(x) = -1$,

$$E_{\text{out}} = P[y = +1] = 1 - \tau$$

3. $x < 0$: $h(x) = -1$,

$$E_{\text{out}} = P[y = +1] = \tau$$

Hence,

$$E_{\text{out}} = \left(\frac{1}{2} - \theta\right)\tau + \theta(1 - \tau) + \frac{1}{2}\tau = \theta(1 - 2\tau) + \tau$$

- $-0.5 < \theta < 0$: by symmetry,

$$E_{\text{out}} = -\theta(1 - 2\tau) + \tau$$

- $\theta < -0.5$: again using symmetry,

$$E_{\text{out}} = \frac{1}{2}$$

For all cases,

$$E_{\text{out}} = \min(|\theta|, 0.5)(1 - 2\tau) + \tau$$

The answer is [d].

13.

```
import numpy as np

def Ein(s, x, y, theta):
    if theta == float('-inf'):
        return np.count_nonzero(y == -s) / len(y)
    h = np.sign(s * (x - theta))
    return np.sum(h != y) / len(y)

def Eout(s, theta):
    return min(abs(theta), 0.5) if s == 1 else 1 - min(abs(theta), 0.5)

def DecisionStump():
    x = np.sort(np.random.uniform(low=-0.5, high=0.5, size=2))
    y = np.sign(x)

    thetas = [float('-inf')] + [(x[i] + x[i + 1]) / 2 for i in range(len(x) - 1)]

    minE = float('inf')
    minS = None
    minTheta = None

    for t in thetas:
        for s in [-1, 1]:
            E = Ein(s, x, y, t)
            if E < minE:
                minE = E
```

```

        minS = s
        minTheta = t
    elif E == minE and s * t < minS * minTheta:
        minE = E
        minS = s
        minTheta = t

    return Eout(minS, minTheta) - minE

mean = sum(DecisionStump() for i in range(10000)) / 10000

print(mean) # 0.29376981584936057

```

The answer is [b].

14.

```

import numpy as np

def Ein(s, x, y, theta):
    if theta == float('-inf'):
        return np.count_nonzero(y == -s) / len(y)
    h = np.sign(s * (x - theta))
    return np.sum(h != y) / len(y)

def Eout(s, theta):
    return min(abs(theta), 0.5) if s == 1 else 1 - min(abs(theta), 0.5)

def DecisionStump():
    x = np.sort(np.random.uniform(low=-0.5, high=0.5, size=128))
    y = np.sign(x)

    thetas = [float('-inf')] + [(x[i] + x[i + 1]) / 2 for i in range(len(x) - 1)]

    minE = float('inf')
    minS = None
    minTheta = None

    for t in thetas:
        for s in [-1, 1]:
            E = Ein(s, x, y, t)
            if E < minE:
                minE = E
                minS = s

```

```

        minTheta = t
    elif E == minE and s * t < minS * minTheta:
        minE = E
        minS = s
        minTheta = t

    return Eout(minS, minTheta) - minE

mean = sum(DecisionStump() for i in range(10000)) / 10000

print(mean) # 0.003907311865304553

```

The answer is [b].

15.

```

import numpy as np

tau = 0.2

def Ein(s, x, y, theta):
    if theta == float('-inf'):
        return np.count_nonzero(y == -s) / len(y)
    h = np.sign(s * (x - theta))
    return np.sum(h != y) / len(y)

def Eout(s, theta):
    return min(abs(theta), 0.5) * (1 - 2 * tau) + tau if s == 1 else 1 -
    min(abs(theta), 0.5) * (1 - 2 * tau) + tau

def DecisionStump():
    x = np.sort(np.random.uniform(low=-0.5, high=0.5, size=2))
    y = np.sign(x) * np.random.choice([1, -1], size=len(x), p=[1 - tau, tau])

    thetas = [float('-inf')] + [(x[i] + x[i + 1]) / 2 for i in range(len(x) - 1)]

    minE = float('inf')
    minS = None
    minTheta = None

    for t in thetas:
        for s in [-1, 1]:
            E = Ein(s, x, y, t)
            if E < minE:
                minE = E

```



```

        minS = s
        minTheta = t
    elif E == minE and s * t < minS * minTheta:
        minE = E
        minS = s
        minTheta = t

    return Eout(minS, minTheta) - minE

mean = sum(DecisionStump() for i in range(10000)) / 10000

print(mean) # 0.5662640455969353

```

The answer is [d].

16.

```

import numpy as np

tau = 0.2

def Ein(s, x, y, theta):
    if theta == float('-inf'):
        return np.count_nonzero(y == -s) / len(y)
    h = np.sign(s * (x - theta))
    return np.sum(h != y) / len(y)

def Eout(s, theta):
    return min(abs(theta), 0.5) * (1 - 2 * tau) + tau if s == 1 else 1 -
    min(abs(theta), 0.5) * (1 - 2 * tau) + tau

def DecisionStump():
    x = np.sort(np.random.uniform(low=-0.5, high=0.5, size=128))
    y = np.sign(x) * np.random.choice([1, -1], size=len(x), p=[1 - tau, tau])

    thetas = [float('-inf')] + [(x[i] + x[i + 1]) / 2 for i in range(len(x) - 1)]

    minE = float('inf')
    minS = None
    minTheta = None

    for t in thetas:
        for s in [-1, 1]:
            E = Ein(s, x, y, t)
            if E < minE:

```

```

        minE = E
        minS = s
        minTheta = t
    elif E == minE and s * t < minS * minTheta:
        minE = E
        minS = s
        minTheta = t

    return Eout(minS, minTheta) - minE

mean = sum(DecisionStump() for i in range(10000)) / 10000

print(mean) # 0.013888179740139808

```

The answer is [b].

17.

```

import numpy as np

def Ein(s, x, y, theta):
    if theta == float('-inf'):
        return np.count_nonzero(y == -s) / len(y)
    h = np.sign(s * (x - theta))
    return np.sum(h != y) / len(y)

def DecisionStump(x, y):
    thetas = [float('-inf')] + [(x[i] + x[i + 1]) / 2 for i in range(len(x) - 1)]

    minE = float('inf')
    minS = None
    minTheta = None

    for t in thetas:
        for s in [-1, 1]:
            E = Ein(s, x, y, t)
            if E < minE:
                minE = E
                minS = s
                minTheta = t
            elif E == minE and s * t < minS * minTheta:
                minE = E
                minS = s
                minTheta = t

    return minE

```

```

data = np.loadtxt("hw2_train.dat")

x = np.transpose(data)[: -1]
y = np.transpose(data)[-1]

print(np.min([DecisionStump(xi, y) for xi in x])) # 0.026041666666666668

```

The answer is [c].

18.

```

import numpy as np

def Ein(s, x, y, theta):
    if theta == float('-inf'):
        return np.count_nonzero(y == -s) / len(y)
    h = np.sign(s * (x - theta))
    return np.sum(h != y) / len(y)

def DecisionStump(x, y):
    thetas = [float('-inf')] + [(x[i] + x[i + 1]) / 2 for i in range(len(x) - 1)]

    minE = float('inf')
    minS = None
    minTheta = None

    for t in thetas:
        for s in [-1, 1]:
            E = Ein(s, x, y, t)
            if E < minE:
                minE = E
                minS = s
                minTheta = t
            elif E == minE and s * t < minS * minTheta:
                minE = E
                minS = s
                minTheta = t

    return [minE, minS, minTheta]

data = np.loadtxt("hw2_train.dat")
test = np.loadtxt("hw2_test.dat")

x = np.transpose(data)[: -1]

```

```

y = np.transpose(data)[-1]
g = np.array([DecisionStump(xi, y) for xi in x])
k = np.argmin(g[:, 0])

Eout = Ein(g[k][1], np.transpose(test)[k], np.transpose(test)[-1], g[k][2])
print(Eout) # 0.078125

```

The answer is [e].

19.

```

import numpy as np

def Ein(s, x, y, theta):
    if theta == float('-inf'):
        return np.count_nonzero(y == -s) / len(y)
    h = np.sign(s * (x - theta))
    return np.sum(h != y) / len(y)

def DecisionStump(x, y):
    thetas = [float('-inf')] + [(x[i] + x[i + 1]) / 2 for i in range(len(x) - 1)]

    minE = float('inf')
    minS = None
    minTheta = None

    for t in thetas:
        for s in [-1, 1]:
            E = Ein(s, x, y, t)
            if E < minE:
                minE = E
                minS = s
                minTheta = t
            elif E == minE and s * t < minS * minTheta:
                minE = E
                minS = s
                minTheta = t

    return [minE, minS, minTheta]

data = np.loadtxt("hw2_train.dat")
test = np.loadtxt("hw2_test.dat")

x = np.transpose(data)[: -1]
y = np.transpose(data)[-1]
g = np.array([DecisionStump(xi, y) for xi in x])

```

```
print(np.max(g[:, 0]) - np.min(g[:, 0])) # 0.3020833333333333
```

The answer is [d].

20.

```
import numpy as np

def Ein(s, x, y, theta):
    if theta == float('-inf'):
        return np.count_nonzero(y == -s) / len(y)
    h = np.sign(s * (x - theta))
    return np.sum(h != y) / len(y)

def DecisionStump(x, y):
    thetas = [float('-inf')] + [(x[i] + x[i + 1]) / 2 for i in range(len(x) - 1)]

    minE = float('inf')
    minS = None
    minTheta = None

    for t in thetas:
        for s in [-1, 1]:
            E = Ein(s, x, y, t)
            if E < minE:
                minE = E
                minS = s
                minTheta = t
            elif E == minE and s * t < minS * minTheta:
                minE = E
                minS = s
                minTheta = t

    return [minE, minS, minTheta]

data = np.loadtxt("hw2_train.dat")
test = np.loadtxt("hw2_test.dat")

x = np.transpose(data)[: -1]
y = np.transpose(data)[-1]
g = np.array([DecisionStump(xi, y) for xi in x])
b = np.argmax(g[:, 0])
k = np.argmin(g[:, 0])

Eoutb = Ein(g[b][1], np.transpose(test)[b], np.transpose(test)[-1], g[b][2])
```

```
Eout = Ein(g[k][1], np.transpose(test)[k], np.transpose(test)[-1], g[k][2])
```

```
print(Eoutb - Eout) # 0.34375
```

The answer is [b].